

# Project Documentation

Power Plant Energy Output Prediction  
Using Artificial Neural Networks (PyTorch)

*ANN Regression | Deep Learning | Combined Cycle Power Plant*

---

# 1. Introduction

---

## 1.1 Problem Statement

Predict the net hourly electrical energy output (PE) of a Combined Cycle Power Plant (CCPP) based on ambient environmental conditions. This is a regression task where the goal is to estimate a continuous target variable.

## 1.2 Objective

Build and train an Artificial Neural Network (ANN) using PyTorch to accurately predict the energy output given four input features: Temperature, Exhaust Vacuum, Ambient Pressure, and Relative Humidity.

## 1.3 Tools & Technologies

Tool	Purpose
Python 3.x	Programming language
PyTorch	Deep learning framework
Pandas	Data loading & manipulation
NumPy	Numerical operations
scikit-learn	Preprocessing & evaluation metrics
Matplotlib	Visualization
Jupyter Notebook	Development environment

# 2. Dataset Description

---

**File:** powerplant\_data.csv

**Records:** 9,568 observations

**Source:** Combined Cycle Power Plant dataset

## 2.1 Feature Details

Column	Full Name	Unit	Role
AT	Ambient Temperature	°C	Input Feature
V	Exhaust Vacuum	cm Hg	Input Feature
AP	Ambient Pressure	millibar	Input Feature
RH	Relative Humidity	%	Input Feature

PE	Produced Energy (Net hourly electrical energy output)	MWh	Target Variable
----	---	-----	-----------------

## 2.2 Data Quality

- Missing values: None (verified with `df.isnull().sum()`)
- Data types: All columns are numeric (`float64`)

## 3. Data Preprocessing

### 3.1 Feature-Target Split

```
X = df.drop("PE", axis=1)      # Features: AT, V, AP, RH
y = df["PE"]                   # Target: PE
```

### 3.2 Train-Test Split

```
x_train, x_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

- Training set: 80% (~7,654 samples)
- Test set: 20% (~1,914 samples)
- Random state: 42 (for reproducibility)

### 3.3 Feature Scaling

Standard scaling (zero mean, unit variance) is applied using `StandardScaler`:

```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)      # fit on train only
X_test_scaled = scaler.transform(X_test)            # transform test using train stats
```

**Why StandardScaler?** Neural networks converge faster and perform better when input features are on a similar scale. `StandardScaler` normalizes each feature to have `mean=0` and `std=1`.

### 3.4 Tensor Conversion & DataLoader

```
X_train_tensor = torch.tensor(X_train_scaled, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train.values, dtype=torch.float32).view(-1, 1)

train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
```

- Data is converted from NumPy/Pandas to PyTorch tensors
- Target is reshaped to (n, 1) for compatibility with the model output
- DataLoader handles batching (batch size = 32) and shuffling

## 4. Model Architecture

### 4.1 Network Design

```
?????????????????????????????????????  
?      Input Layer      ?  
?      (4 neurons)      ?  
?      AT, V, AP, RH      ?  
?????????????????????????????????  
?  
?????????????????????????????????  
?      Hidden Layer 1      ?  
?      Linear(4, 6) + ReLU      ?  
?      (6 neurons)      ?  
?????????????????????????????????  
?  
?????????????????????????????????  
?      Hidden Layer 2      ?  
?      Linear(6, 6) + ReLU      ?  
?      (6 neurons)      ?  
?????????????????????????????????  
?  
?????????????????????????????????  
?      Output Layer      ?  
?      Linear(6, 1)      ?  
?      (1 neuron ? PE)      ?  
?????????????????????????????????
```

### 4.2 Implementation

```

class ANN(nn.Module):
    def __init__(self):
        super(ANN, self).__init__()
        self.model = nn.Sequential(
            nn.Linear(4, 6),      # Hidden Layer 1
            nn.ReLU(),
            nn.Linear(6, 6),      # Hidden Layer 2
            nn.ReLU(),
            nn.Linear(6, 1),      # Output Layer
        )

    def forward(self, x):
        return self.model(x)

```

## 4.3 Design Choices

Choice	Value	Rationale
Hidden layers	2	Sufficient for learning non-linear patterns in tabular data.
Neurons per hidden layer	6	Compact architecture to avoid overfitting on a small dataset.
Activation function	ReLU	Standard choice; avoids vanishing gradient problem.
Output activation	None (linear)	Regression task requires unbounded continuous output.

## 4.4 Total Parameters

Layer	Parameters
Linear(4, 6)	$4 \times 6 + 6 = **30**$
Linear(6, 6)	$6 \times 6 + 6 = **42**$
Linear(6, 1)	$6 \times 1 + 1 = **7**$
**Total**	<b>**79**</b>

## 5. Training

### 5.1 Configuration

Hyperparameter	Value

Loss Function	MSELoss (Mean Squared Error)
Optimizer	Adam (default lr=0.001)
Epochs	100
Batch Size	32

## 5.2 Training Loop

For each epoch:

- Training phase (model.train()):
- Iterate over mini-batches from train\_loader
- Forward pass ? compute predictions
- Compute MSE loss
- Backward pass ? compute gradients
- Update weights with Adam optimizer
- Accumulate batch loss for epoch-level tracking
- Validation phase (model.eval() + torch.no\_grad()):
- Iterate over test\_loader without gradient computation
- Compute MSE loss on test data
- Model checkpointing:
- If validation loss improves, save model weights to best\_model.pt

## 5.3 Training Output

Each epoch prints:

```
epoch $1/100 ==> train loss = X.XXX & valid = X.XXX
```

# 6. Evaluation

## 6.1 Best Model Loading

```
model.load_state_dict(torch.load("best_model.pt"))
```

The model with the lowest validation loss across all epochs is loaded for final evaluation.

## 6.2 Metrics

Metric	Description	Formula

**MSE** (Mean Squared Error)	Average squared difference between predicted and actual values
**R <sup>2</sup> Score**	Proportion of variance in the target explained

## 6.3 Evaluation Code

```
model.eval()
with torch.no_grad():
    train_pred = model(X_train_tensor)
    test_pred = model(X_test_tensor)
    train_mse = criterion(train_pred, y_train_tensor)
    test_mse = criterion(test_pred, y_test_tensor)

    r2 = r2_score(y_test, test_pred)
```

## 7. Visualization

### 7.1 Loss Curves

A line plot comparing training loss and validation loss across all 100 epochs is generated using Matplotlib. This helps diagnose:

- Overfitting ? validation loss diverges from training loss
- Underfitting ? both losses remain high
- Good fit ? both losses converge and stabilize

### 7.2 Predictions vs Actuals

A side-by-side DataFrame of predicted and actual energy values is created for qualitative inspection:

```
pd.concat([predicted_df, actual_df], axis=1)
```

## 8. Project File Structure

```
powerplant/
?
?? ANN_Regression.ipynb      # Main Jupyter Notebook (full pipeline)
?? powerplant_data.csv        # Dataset (9,568 records, 5 columns)
?? best_model.pt              # Saved PyTorch model weights (best validation loss)
?? README.md                  # Project overview
?? PROJECT_DOCUMENTATION.md  # This file ? detailed documentation
```

## 9. How to Run

- Clone/download the project folder.
- Install dependencies:

```
pip install torch pandas numpy scikit-learn matplotlib
```

- Open ANN\_Regression.ipynb in Jupyter Notebook or VS Code.
- Run all cells sequentially. The notebook will:
  - Load and preprocess data
  - Define and train the ANN for 100 epochs
  - Save the best model to best\_model.pt
  - Display loss curves and evaluation metrics

## 10. Key Concepts Used

Concept	Description
**ANN (Artificial Neural Network)**	A feed-forward neural network composed of interconnected layers of neurons.
**Regression**	Predicting a continuous numerical value (energy output in MW)
**Backpropagation**	Algorithm to compute gradients of the loss w.r.t. model weights
**Adam Optimizer**	Adaptive learning rate optimizer combining momentum and RMSProp
**MSE Loss**	Penalizes large prediction errors quadratically
**StandardScaler**	Normalizes features to zero mean and unit variance
**Mini-batch Gradient Descent**	Updates weights using small subsets (batches) of data

**Model Checkpointing**	Saving the best model during training to prevent losing optimal weights
**Train/Validation Split**	Separating data to detect overfitting and evaluate generalization

## 11. Potential Improvements

- Add dropout layers to reduce overfitting
- Experiment with deeper/wider architectures
- Use learning rate scheduling (e.g., ReduceLROnPlateau)
- Apply cross-validation for more robust evaluation
- Add early stopping to avoid unnecessary training epochs
- Hyperparameter tuning with Optuna or GridSearch
- Add RMSE and MAE as additional evaluation metrics